

Introduction to Message-Passing Interface

Luca Tornatore - I.N.A.F.



DATA SCIENCE &
ARTIFICIAL INTELLIGENCE



SCIENTIFIC &
DATA-INTENSIVE COMPUTING

2024-2025 @ Università di Trieste

Outline




Intro to
MPI



Point-to-Point
Communications



Collective
Communications



few JEDI things

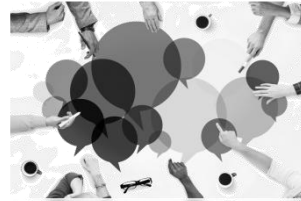
Outline



Intro to
MPI



Point-to-Point
Communications



Collective
Communications



few JEDI things



As we have seen in the “Introduction to parallelism”, the Message-Passing Interface implements a standard that refers to the distributed-memory paradigm of parallel programming.

Hence: MPI implements the *distributed memory* parasigm, and deals with how to *move data* among separate address spaces.



Why MPI

- The MPI *tasks* are **separate processes**, each having its own address space
 - the programmer manages the memory as it was a serial code
- Tasks communicate for
 - Synchronization
 - Data movement
- The MPI tasks exchange data *only* by **collaborative explicit messages** (*)
 - the programmer is in charge of making available memory regions
 - every data distribution among tasks also happens via explicit messages, also managed by the programmer

(*) Actually, since MPI-2, but at mature level from MPI-3, direct memory access is possible among tasks. These are *called one-sided communications* exactly because they do not require the co-operation of the tasks. We'll see some details in week2



Let's specify that MPI

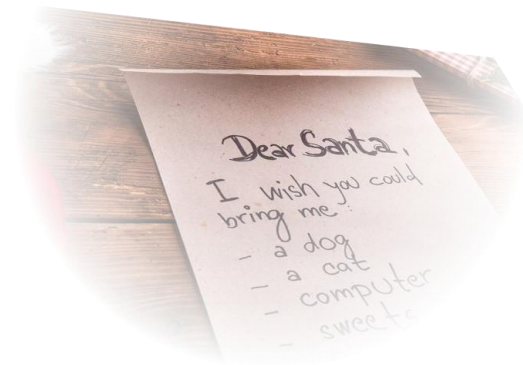
- It is “a book”, i.e. the specification of a standard, *not* an implementation of it.
- Defines **a library, not a language**. All operations depends on the execution of routines.
- It defines the API and the behaviour of the functions; your FORTRAN / C / C++ program is compiled by your compiler after a wrapping by the mpicc program, and then linked to that library.



| Let's invent MPI

If you have to “send a message”,
let's say a mail..

What do you need ?






| Let's invent MPI

If you have to “send a message”, let's say a mail..
What do you need ?

1. a paper **sheet**
2. an **envelop**
3. the name of the **receiver**
4. the **sender's** name, to get an answer
(or a receipt acknowledgment)
5. and, of course a **postal service**
 - a) postal office
 - b) postal network
 - c) postpersons





| Let's invent MPI

Well, that is, in general, exactly how MPI works..

```
char *data_region = (char*)malloc( ... );
```

```
// do something on the data region
```

```
SEND_A_MESSAGE( data_region, how_many_sheets, to_who, in_what_country );
```

the **paper sheet**

how long

the receiver

*his "name" is also the
address if there are only
unique names*

*The community, or
"framework" we are
referring to.
You may belong to
different communities,
right?*




| Let's invent MPI

When you receive a letter, however, in order to get it at the right place - on your desk, in your files, ... - you have to perform some actions.

It may be that you just get a notification, and you have to walk to the postal office, or the you simply have to check your inbox and pick up what is inside.

In short, to receive a letter, you must be collaborative and perform some tasks:



```
RECV_A_MESSAGE( where_to_put_it,  
                how_long,  
                from_who,  
                in_what_country );
```

← of course you need some
room for it

← and you need to know how long it is

← better to know who is the sender

← and the "framework" he belongs to



| Let's invent MPI

So, we expect that there must be **point-to-point**

- a way to **send** messages
- a corresponding way to **receive** messages
- a way to know “the **names**” in “the **framework**”

what if the
receiver *does*
not expect a
message?





| Let's invent MPI

Since sometimes we need to broadcast messages, it would be nice if

- there was a **broadcasting** (**one-to-many**)

« doomsday will be at
this time at this place
rsvp »

Since we may also in the need of receiving answers, it would also be nice if

- there was a **collection** mechanism (**many-to-one**) mechanism
- the answer/collection was possible **many-to-many** (try to organize a meeting without that..)

« I won't be able to attend
sincerely»



| Let's invent MPI

We're lucky, because MPI designers thought the same.

And that is exactly what we are covering in the next days:
P2P communications and collective communications.

They exist in several different flavours, of which we'll explore
definition and usage.



| The very basic

However, we need to start with MPI before using its routines.
And that is the most simple MPI code you may ever write:

```
#include <mpi.h>
int main ( int argc, char **argv )
{
    MPI_Init( &argc, &argv );

    MPI_Finalize();
    return 0;
}
```

just try it:

```
mpicc verybasic.c -o verybasic
mpirun -np $NUM ./verybasic
```




The very basic

raise your hand if you have any doubts about these two things..



```
#include <mpi.h>
```

```
int main ( int argc, char **argv )
```

```
{
```

```
    MPI_Init( &argc, &argv );
```

```
    MPI_Finalize();
```

```
    return 0;
```

```
}
```

All MPI routine calls begin with **MPI_**

You always need to **initialize MPI** first.
Until we don't do that, you'll just have a bunch of processes doing the same thing, and you probably failing at the first MPI call.

Best practice: always finalize your environment

Best practice: always explicitly return, and return with a value





| The very basic

Better to get used to **MPI_Init_thread()**

```
int main ( int argc, char **argv )
{
    int mpi_provided_thread_level;
    MPI_Init_thread( &argc, &argv, REQUIRED_LEVEL, &mpi_provided_thread_level );

    if ( mpi_provided_thread_level < REQUIRED_LEVEL ) {
        ..manage the situation.. }

    ...
}
```

where **REQUIRED_LEVEL** can be:

MPI_THREAD_SINGLE	Only one thread will execute.
MPI_THREAD_FUNNELED	Only the thread that called MPI_Init_thread will make MPI calls.
MPI_THREAD_SERIALIZED	Only one thread will MPI library calls at one time.
MPI_THREAD_MULTIPLE	Multiple threads may call MPI at once with no restrictions.



| The very basic: init and end

The second most simple MPI code you may ever write, is, obviously “hello MPI world”

```
#include <mpi.h>

int main ( int argc, char **argv )
{
    int provided_thread_level;
    MPI_Init_thread( &argc, &argv, MPI_THREAD_SINGLE, &provided_thread_level );

    printf(“hello MPI world\n”);

    MPI_Finalize();
    return 0;
}
```

Let's try that and see what happens



| The very basic: who am I?

Definitely something was missing, and the following is the real “hello MPI world” thing

```
int main ( int argc, char **argv )
{
    int Ntasks, Myrank;
    int provided_thread_level;
    MPI_Init_thread( &argc, &argv, MPI_THREAD_SINGLE, &provided_thread_level );

    MPI_Comm_size ( MPI_COMM_WORLD, &Ntasks ); # OF HOW MANY TASK ARE AT WORK
    MPI_Comm_rank ( MPI_COMM_WORLD, &Myrank ); WHO IS THE TASK THAT ARE EXECUTING (0→n-1)
    printf(“hello there MPI world from task %d out of %d\n”,
           Myrank, Ntasks );

    MPI_Finalize();
    return 0;
}
```

[see helloworld.c](#)



The very basic: who am I ?

Definitely something was missing, and the following is the real “hello MPI world” thing

```
int main ( int argc, char **argv )
{
    int Ntasks, Myrank;
    int provided_thread_level;
    MPI_Init_thread( &argc, &argv, MPI_THREAD_SINGLE, &provided_thread_level );

    MPI_Comm_size ( MPI_COMM_WORLD, &Ntasks );
    MPI_Comm_rank ( MPI_COMM_WORLD, &Myrank );
    printf(“hello there MPI world from task %d out of %d\n”,
           Myrank, Ntasks );

    MPI_Finalize();
    return 0;
}
```

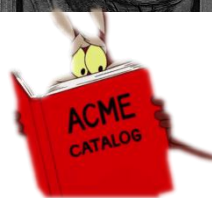
Gets how many tasks are in the communicator

Gets the rank of the calling process in the communicator



| The very basic: an hint

That is the moment for an **hint**..



There exist a wonderful thing, invented by smart people for people like me, which can not remember too complicate stuff..

the **man page**!

Try `man MPI_$the_routine_you_want_to_discover`



| The very basic: the communicators

Back to here... what is “a communicator” ?

```
int main ( int argc, char **argv )  
{
```

```
    int Ntasks, Myrank;
```

```
    MPI_Init_thread( &argc, &argv, MPI_THREAD_SERIALIZED );
```

Gets how many tasks are in the **communicator**

```
    MPI_Comm_size ( &Ntasks, MPI_COMM_WORLD );
```

```
    MPI_Comm_rank ( &Myrank, MPI_COMM_WORLD );
```

```
    printf(“hello there MPI world from task %d out of %d\n”,  
           Myrank, Ntasks );
```

```
    MPI_Finalize();
```



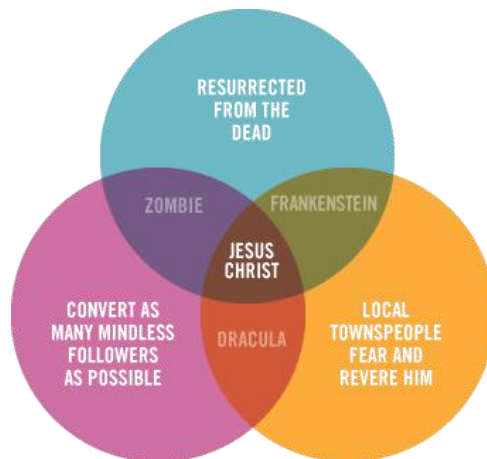
The very basic: the communicators

Communicators and **groups** are a very central concepts in MPI.

Tasks can form **groups**:

A task can be in more than one **group** at a time

The same group can be in different situations



**You and your
colleagues at work
and at a rave**

(didn't put a picture for privacy reasons)



| The very basic: the communicators

Communicators and **groups** are a very central concepts in MPI.

The same group can be
in different situations

**You and your
colleagues at work
and at a rave**

(didn't put a picture for privacy
reasons)

A **COMMUNICATOR** is the combination of a group and its “context”.

You can build as many groups as you want, and they may or not have a communicator.

However, if you want to communicate among tasks in a group, you need a communicator,

(for instance, to assign a rank to every task in the group; for instance your name may be “Gandalf” at work, but to address you at the party we need to call you “Nuanda”)



| The very basic: the communicators

Communicators and **groups** are a very central concepts in MPI.

A **COMMUNICATOR** is the combination of a group and its “context”.

- This functionality offers the capability of isolating communication between application modules with an effective “sandbox” for different contexts. For instance, a parallel library and your application will use internally their own communicator, separating contexts.
- By creating groups of MPI processes, that may or not overlap with each other, it is possible to
 - separate contexts within different modules of the same application (useful or even *advisable*)
 - express multiple levels of parallelism



| The very basic: the communicators

Communicators and **groups** are a very central concepts in MPI.

A **COMMUNICATOR** is the combination of a group and its “context”.

- **MPI_COMM_WORLD**
is the default communicator available right after the call to MPI_Init.
Its group contains *all* the tasks started by your job.
- **MPI_COMM_NULL**
signals an invalid / non existent communicator
- **MPI_COMM_SELF**
contains only the process itself
- **MPI_GROUP_NULL**
signals an invalid / non existent group



BEST PRACTICE



always create a separated “context” for the application you’re writing

```
#include <mpi.h>
```

```
int main ( int argc, char **argv )  
{
```

```
    int Myrank, Ntasks;  
    int mpi_provided_thread_level;  
    MPI_Comm myCOMM_WORLD;
```

```
    MPI_Init_thread ( &argc, &argv, MPI_THREAD_SINGLE, &mpi_provided_thread_level);  
    MPI_Comm_dup ( MPI_COMM_WORLD, &myCOMM_WORLD );
```

```
    MPI_Comm_size( &Ntasks, myCOMM_WORLD);  
    MPI_Comm_rank( &Myrank, myCOMM_WORLD);
```

```
    ...
```




The very basic: the communicators

BEST PRACTICE



always create a separated “context” for the application you’re writing

Quoting Victor Eijkhout:

« Imagine you’re writing a library, and your library makes MPI calls. Now imagine that some `Isend` and `Irecv` calls are done in a library routine that the user calls, with the `Wait` calls in another routine that the library calls.

Since user code is active in between the library doing `Isend/Irecv` and the library doing wait, it is possible for the user to catch the library `Isend/Irecv` calls, or conversely for the user to start `Isend/Irecv` calls and the library to catch them.

You prevent such mishaps by letting the library use a duplicate communicator on the same group of processes. The communicator in effect becomes a label on the messages. »